

# HW4\_report

April 4, 2016

## 1 CS189: Introduction to Machine Learning

### 1.1 Problem 1: Ridge Regression

In this problem we will return to predicting the median home value in a given Census area by extending linear regression. The data is in `housing_data.mat` and it comes from <ftp://rcom.univie.ac.at/mirrors/lib.stat.cmu.edu/datasets/.index.html>. There are only 8 features for each data point; you can read about the features in `housing_data_source.txt`.

#### 1.1.1 1a)

In order to find the optimizer of the loss function  $J(w, \alpha)$ , we need to differentiate the loss function  $J(w, \alpha)$  and find the local minima. Before differentiating the loss function  $J(w, \alpha)$ , let us expand the mathematical expression to facilitate calculations.

$$\begin{aligned} J(w, \alpha) &= (Xw + \alpha \mathbf{1} - y)^\top (Xw + \alpha \mathbf{1} - y) + \lambda w^\top w \\ &= (w^\top X^\top + \alpha \mathbf{1}^\top - y^\top)(Xw + \alpha \mathbf{1} - y) + \lambda w^\top w \\ &= w^\top X^\top Xw + \alpha w^\top X^\top \mathbf{1} - w^\top X^\top y + \alpha \mathbf{1}^\top Xw + \alpha^2 n - \alpha \mathbf{1}^\top y + -y^\top Xw - y^\top \alpha \mathbf{1} + y^\top y + \lambda w^\top w \\ &= w^\top X^\top Xw - w^\top X^\top y + \alpha n - \alpha \mathbf{1}^\top y + -y^\top Xw - y^\top \alpha \mathbf{1} + y^\top y + \lambda w^\top w \quad (\because X^\top \mathbf{1} = \mathbf{1}^\top X = 0) \end{aligned}$$

If we differentiate the loss function  $J(w, \alpha)$  by  $w$ , then

$$\frac{\partial J}{\partial w} = 2X^\top Xw - 2X^\top y + 2\lambda w.$$

Therefore,

$$\frac{\partial J}{\partial w} = 2X^\top Xw - 2X^\top y + 2\lambda w = 0$$

gives  $\hat{w} = (X^\top X + \lambda I)^{-1} X^\top y$ .

If we differentiate the loss function  $J(w, \alpha)$  by  $\alpha$ , then we get

$$\frac{\partial J}{\partial \alpha} = 2n - 2\mathbf{1}^\top y.$$

In order for  $\frac{\partial J}{\partial \alpha}$  to be 0, it should be that  $\hat{\alpha} = \bar{y}$ .

#### 1.1.2 1b)

```
In [2]: import scipy.io as sio
import numpy as np
```

```
housing_data = sio.loadmat('./housing_dataset/housing_data.mat')
```

```

In [28]: housing_xtrain = housing_data['Xtrain']
         housing_ytrain = housing_data['Ytrain'][:,0]
         housing_xvalid = housing_data['Xvalidate']
         housing_yvalid = housing_data['Yvalidate'][:,0]

In [46]: center_housing_xtrain.shape

Out[46]: (19440, 8)

In [19]: # center the feature values of the training and validation data.

         center_housing_xtrain = housing_xtrain - np.expand_dims(np.mean(housing_xtrain, axis = 0), axis=1)
         center_housing_xvalid = housing_xvalid - np.expand_dims(np.mean(housing_xvalid, axis = 0), axis=1)

In [23]: np.sum(center_housing_xvalid, axis=0)

Out[23]: array([ 4.19220214e-13, -2.70006240e-13, -9.45874490e-11,
                 -1.90993887e-11, -8.36735126e-11,  2.91038305e-11,
                 2.77111667e-13, -2.55795385e-13])

In [5]: from numpy.linalg import inv

In [64]: # Implement ridge regression for (i)
         def ridge_regression (X, y, lamda):
             w = np.dot(inv(np.dot(np.transpose(X), X) + (lamda * np.identity(8))),
                         np.dot(np.transpose(X), y))
             alpha = np.mean(y)
             return [w, alpha]

In [6]: # Implement 10-fold Cross-Validation, using the ridge_regression implemented above.

         from sklearn.cross_validation import KFold
         kf = KFold(19440, n_folds=10)

         def tenfold_cv(lamda):
             RSS = 0
             for train_index, test_index in kf:
                 X_train, X_test = center_housing_xtrain[train_index, :], center_housing_xtrain[test_index, :]
                 y_train, y_test = housing_ytrain[train_index], housing_ytrain[test_index]
                 ridge = ridge_regression(X_train, y_train, lamda)
                 w = ridge[0]
                 alpha = ridge[1]
                 one = np.linspace(1,1,len(y_test))
                 temp = (np.dot(X_test, w) + (alpha * np.transpose(one)) - y_test )
                 RSS = RSS + np.dot(np.transpose(temp), temp)
             return RSS/10

In [88]: tenfold_cv(1e-12)

Out[88]: 9425875042980.4492

In [89]: tenfold_cv(1e-5)

Out[89]: 9425875042967.5977

```

```

In [90]: tenfold_cv(1e-4)
Out[90]: 9425875042851.9395
In [91]: tenfold_cv(1e-3)
Out[91]: 9425875041695.3809
In [92]: tenfold_cv(1e-2)
Out[92]: 9425875030137.1055
In [93]: tenfold_cv(1e-1)
Out[93]: 9425874915283.6348
In [94]: tenfold_cv(1)
Out[94]: 9425873839646.3105
In [95]: tenfold_cv(10)
Out[95]: 9425870341696.4805
In [96]: tenfold_cv(100)
Out[96]: 9426530937596.7383
In [97]: tenfold_cv(1000)
Out[97]: 9480399010571.4824
In [98]: tenfold_cv(5)
Out[98]: 9425870659168.0137
In [99]: tenfold_cv(20)
Out[99]: 9425881801097.7461
In [100]: tenfold_cv(12)
Out[100]: 9425871347405.1973
In [101]: tenfold_cv(8)
Out[101]: 9425869982405.9648
In [102]: tenfold_cv(7)
Out[102]: 9425870045690.5527
In [103]: tenfold_cv(9)
Out[103]: 9425870081144.416

```

Average RSS on different tuned lambda valuse is given as follows:

Lambda	average RSS
1e-12	9425875042980.4492
1e-5	9425875042967.5977
1e-4	9425875042851.9395
1e-3	9425875041695.3809
1e-2	9425875030137.1055
1e-1	9425874915283.6348
1	9425873839646.3105
5	9425870659168.0137
7	9425870045690.5527
8	9425869982405.9648
9	9425870081144.416
10	9425870341696.4805
12	9425871347405.1973
20	9425881801097.7461
100	9426530937596.7383
1000	9480399010571.4824

It is when  $\lambda = 8$  that ridge regression with least squares achieves the lowest average RSS throughout 10-fold cross validation.

In [109]: *# RSS on validation data with the lambda 8 with the lowest cross-validation error.*

```
ridge = ridge_regression(center_housing_xtrain, housing_ytrain, 8)
w = ridge[0]
alpha = ridge[1]
one = np.linspace(1,1,len(housing_yvalid))
temp = (np.dot(center_housing_xvalid, w) + (alpha * np.transpose(one)) - housing_yvalid )
RSS = np.dot(np.transpose(temp), temp)
```

In [110]: RSS

Out[110]: 5782552884504.541

When we train ridge regression with lambda of 8, RSS on validation data is given as 5782552884504.541. Compared to 5794953797654.9834, RSS of linear regression in HW3, the ridge regression exhibits lower RSS and therefore fits better to the validation data than the linear regression does.

In [111]: w

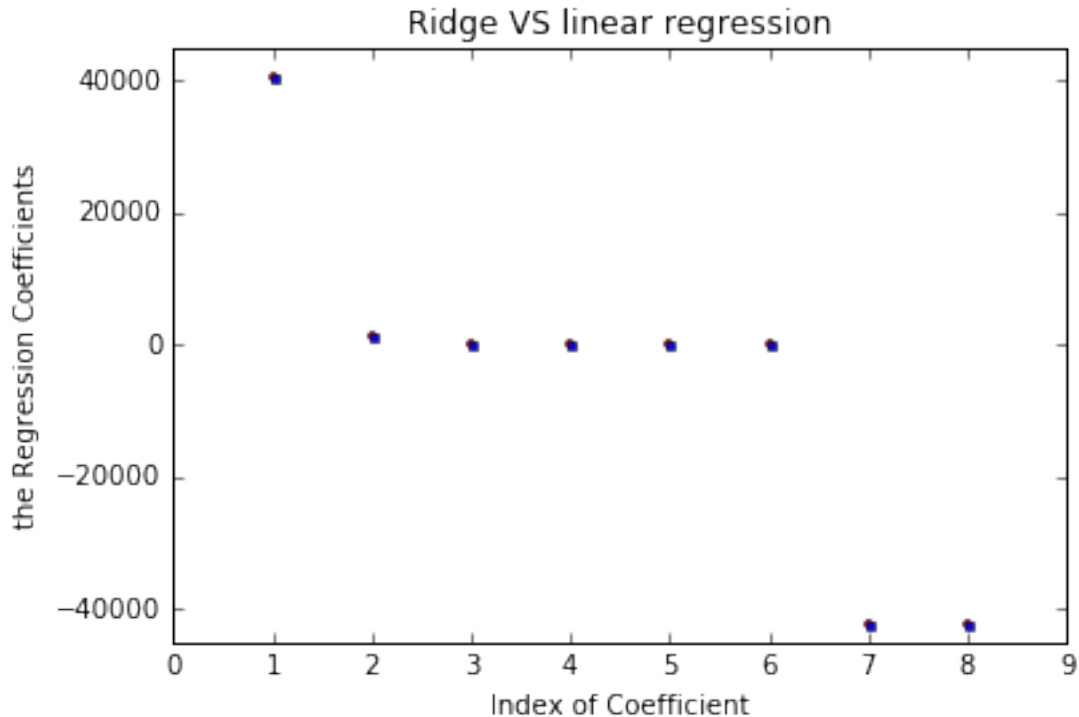
Out[111]: array([ 4.05915089e+04, 1.19668415e+03, -8.50719390e+00,  
1.18261856e+02, -3.77927882e+01, 4.32070327e+01,  
-4.21178034e+04, -4.23914057e+04])

In [115]: beta = np.array([ 4.05879986e+04, 1.19561189e+03, -8.50145688e+00,  
1.18352188e+02, -3.77900280e+01, 4.30562637e+01,  
-4.21794075e+04, -4.24573474e+04])

In [7]: import matplotlib.pyplot as plt  
%matplotlib inline

In [124]: xpos = np.arange(len(w))+ 1  
plt.axis([0,9, -45000, 45000])  
plt.plot(xpos, w, 'ro', ms = 3)

```
plt.plot(xpos, beta, 'bs', ms =3)
plt.xlabel('Index of Coefficient')
plt.ylabel('the Regression Coefficients')
plt.title('Ridge VS linear regression')
plt.show()
```



We plot the regression coefficients of 8 feature in both ridge regression and linear regression in HW3. This plot shows that the maximum absolute value of the regression coefficients is slightly smaller in the ridge regression than in the linear regression, as we could expected. Since we impose penalty on large norm of regression coefficient vector in ridge regression, the ridge regression tends to have smaller absolute value of regression coefficients.

## 1.2 Problem 2 : Logistic Regression

```
In [131]: X = np.array([[0,3,1],[1,3,1],[0,1,1],[1,1,1]])
          y = np.transpose(np.array([1,1,0,0]))
          1/(np.array([1,2,3]) + 1)
```

```
Out[131]: array([ 0.5          ,  0.33333333,  0.25          ])
```

```
In [8]: # Implement the logistic function
def s(gamma):
    return 1/(1+ np.exp(- gamma))

# Implement the logistic loss, aka the cross-entropy loss.
def R(w,X,y):
    temp = np.dot(X,w)
    one = np.linspace(1,1, len(y))
```

```

        return np.dot(np.transpose(1- y),temp) + np.dot(one, np.log(1+ np.exp(-temp)))

        #return - np.dot(np.transpose(y) ,np.log(s(np.dot(X, w)))) - np.dot(np.transpose(1-y), np.l

def batch_gradient(w,X,y,e):
    return w + e* np.dot(np.transpose(y- s(np.dot(X, w))), X)

In [266]: w0 = np.transpose(np.array([-2,1,0]))
In [269]: R(w0, X, y)
Out[269]: 1.9883724141284105
In [219]: mu0 = s(np.dot(X,w0))
           mu0
Out[219]: array([ 0.95257413,  0.73105858,  0.73105858,  0.26894142])
In [220]: w1 = batch_gradient(w0, X, y, 1)
           w1
Out[220]: array([-2.          ,  0.94910188, -0.68363271])
In [221]: mu1 = s(np.dot(X,w1))
           mu1
Out[221]: array([ 0.89693957,  0.54082713,  0.56598026,  0.15000896])
In [222]: w2 = batch_gradient(w1, X, y, 1)
           w2
Out[222]: array([-1.69083609,  1.91981257, -0.83738862])
In [271]: R(w2, X, y)
Out[271]: 1.8546997847922486

```

All things considered, \* the value of  $R(w_0)$  is 1.9883724141284105, \* the value of  $\mu_0$  is (0.95257413, 0.73105858, 0.73105858, 0.26894142), \* the value of  $w_1$  is  $[-2., 0.94910188, -0.68363271]^T$ , \* the value of  $\mu_1$  is (0.89693957, 0.54082713, 0.56598026, 0.15000896), \* the value of  $w_2$  is  $[-1.69083609, 1.91981257, -0.83738862]^T$ , \* the value of  $R(w_2)$  is 1.8546997847922486.

### 1.3 Problem 3: Spam classification using Logistic Regression

The spam dataset given as part of the homework in spam.mat consists of 5172 email messages, from which 32 features have been extracted. Please use the standard features for the first four parts of this problem. In part 5, we are asked to predict the labels of the test set in test.mat and submit the predictions to Kaggle. Feel free to use your own featurizes to boost up your score! One can imagine performing several kinds of preprocessing to this data matrix. Try each of the following separately:

1. Standardize each column to have mean 0 and unit variance.
2. Transform the features using  $X_{ij} \mapsto \log(X_{ij} + 0.1)$ , where the  $X_{ij}$  's are the entries of the design matrix.
3. Binarize the features using  $X_{ij} \mapsto I(X_{ij} > 0)$ .  $I$  denotes an indicator variable.

```
In [21]: # import the data
```

```

spam_data = sio.loadmat('./spam_dataset/spam_data.mat')
sp_test = spam_data['test_data']
sp_xtrain = spam_data['training_data']
sp_ytrain = spam_data['training_labels'][0,]
sp_test = np.hstack((sp_test, np.expand_dims(np.transpose(np.linspace(1,1, sp_test.shape[0])),

```

## Preprocessing

- Standardize each column to have mean 0 and unit variance.

In [22]: *# Standardize each column to have mean 0 and unit variance.*

```
col_mean = np.mean(sp_xtrain, axis=0)
col_std = np.std(sp_xtrain, axis = 0)
```

```
sp_xtrain1 = (sp_xtrain - np.expand_dims(col_mean,axis = 0))/np.expand_dims(col_std, axis = 0)
sp_xtrain1 = np.hstack((sp_xtrain1, np.expand_dims(np.transpose(np.linspace(1,1, sp_xtrain1.sh
```

- Transform the features using  $X_{ij} \mapsto \log(X_{ij} + 0.1)$ , where the  $X_{ij}$  's are the entries of the design matrix.

In [23]: `sp_xtrain2 = np.log(sp_xtrain + 0.1)`

```
sp_xtrain2 = np.hstack((sp_xtrain2, np.expand_dims(np.transpose(np.linspace(1,1, sp_xtrain2.sh
```

- Binarize the features using  $X_{ij} \mapsto I(X_{ij} > 0)$ .  $I$  denotes an indicator variable.

In [24]: `sp_xtrain3 = (sp_xtrain > 0) * 1`

```
sp_xtrain3 = np.hstack((sp_xtrain3, np.expand_dims(np.transpose(np.linspace(1,1, sp_xtrain3.sh
```

**Batch Gradient Descent** We run the batch gradient method with learning rate  $10^{-3}$  for the first preprocessing data.

In [25]: `w_0 = np.transpose(np.zeros(sp_xtrain1.shape[1]))`

In [26]: `def logistic_with_batch_gradient(w,X,y,e,n):`

```
    w_temp = w
    risk = []
    for i in np.arange(n):
        risk_temp = R(w_temp,X,y)
        print ('the risk is ', risk_temp, ': iteration ', (i+1))
        w_temp = batch_gradient(w_temp, X, y, e)
        risk = risk + [risk_temp]
    return [risk, w_temp]
```

In [27]: `empirical_risk1 = logistic_with_batch_gradient(w_0, sp_xtrain1, sp_ytrain, 1e-3, 2000)`

```
the risk is 3584.95721786 : iteration 1
the risk is 2387.22054291 : iteration 2
the risk is 2228.25903423 : iteration 3
the risk is 2202.00609786 : iteration 4
the risk is 2184.67717818 : iteration 5
the risk is 2171.56129252 : iteration 6
the risk is 2161.02106254 : iteration 7
the risk is 2152.25131949 : iteration 8
the risk is 2144.779402 : iteration 9
the risk is 2138.29922082 : iteration 10
the risk is 2132.59971606 : iteration 11
the risk is 2127.52901278 : iteration 12
the risk is 2122.97444877 : iteration 13
the risk is 2118.85047442 : iteration 14
the risk is 2115.09083712 : iteration 15
the risk is 2111.64331447 : iteration 16
```

```

the risk is 2009.00718748 : iteration 1961
the risk is 2009.00448491 : iteration 1962
the risk is 2009.00178461 : iteration 1963
the risk is 2008.99908658 : iteration 1964
the risk is 2008.99639081 : iteration 1965
the risk is 2008.99369729 : iteration 1966
the risk is 2008.99100603 : iteration 1967
the risk is 2008.98831702 : iteration 1968
the risk is 2008.98563026 : iteration 1969
the risk is 2008.98294574 : iteration 1970
the risk is 2008.98026347 : iteration 1971
the risk is 2008.97758344 : iteration 1972
the risk is 2008.97490565 : iteration 1973
the risk is 2008.97223009 : iteration 1974
the risk is 2008.96955677 : iteration 1975
the risk is 2008.96688567 : iteration 1976
the risk is 2008.96421681 : iteration 1977
the risk is 2008.96155017 : iteration 1978
the risk is 2008.95888574 : iteration 1979
the risk is 2008.95622354 : iteration 1980
the risk is 2008.95356356 : iteration 1981
the risk is 2008.95090578 : iteration 1982
the risk is 2008.94825022 : iteration 1983
the risk is 2008.94559687 : iteration 1984
the risk is 2008.94294572 : iteration 1985
the risk is 2008.94029677 : iteration 1986
the risk is 2008.93765003 : iteration 1987
the risk is 2008.93500548 : iteration 1988
the risk is 2008.93236312 : iteration 1989
the risk is 2008.92972296 : iteration 1990
the risk is 2008.92708499 : iteration 1991
the risk is 2008.9244492 : iteration 1992
the risk is 2008.9218156 : iteration 1993
the risk is 2008.91918418 : iteration 1994
the risk is 2008.91655493 : iteration 1995
the risk is 2008.91392786 : iteration 1996
the risk is 2008.91130297 : iteration 1997
the risk is 2008.90868024 : iteration 1998
the risk is 2008.90605969 : iteration 1999
the risk is 2008.9034413 : iteration 2000

```


2000 iterations lower the training risk down to 2008.9034413. We plot the training risk (the empirical risk of the training set) vs. the number of iterations.

```

In [28]: xpos = np.arange(len(empirical_risk1[0]))+ 1
plt.plot(xpos, empirical_risk1[0], 'ro', ms =0.4)
plt.xlabel('Iterations')
plt.ylabel('the Training risk')
plt.title('Batch gradient with the First preprocessing method')
plt.show()

```





We run the batch gradient method with learning rate  $10^5$  for the second preprocessing data.

```
In [29]: empirical_risk2 = logistic_with_batch_gradient(w_0, sp_xtrain2, sp_ytrain, 1e-5, 2000)
```

```
the risk is 3584.95721786 : iteration 1
the risk is 3297.42549882 : iteration 2
the risk is 3062.14542981 : iteration 3
the risk is 3017.58606369 : iteration 4
the risk is 2983.2229002 : iteration 5
the risk is 2954.59019699 : iteration 6
the risk is 2928.11948285 : iteration 7
the risk is 2903.12739095 : iteration 8
the risk is 2879.32569929 : iteration 9
the risk is 2856.59869658 : iteration 10
the risk is 2834.87595794 : iteration 11
the risk is 2814.10140317 : iteration 12
the risk is 2794.22548882 : iteration 13
the risk is 2775.20178215 : iteration 14
the risk is 2756.98658009 : iteration 15
the risk is 2739.53842088 : iteration 16
the risk is 2722.81804465 : iteration 17
the risk is 2706.78824526 : iteration 18
the risk is 2691.41380139 : iteration 19
the risk is 2676.66137714 : iteration 20
the risk is 2662.49943896 : iteration 21
the risk is 2648.8981665 : iteration 22
the risk is 2635.8293684 : iteration 23
the risk is 2623.26639864 : iteration 24
```

```

the risk is 1923.50301295 : iteration 1969
the risk is 1923.46869388 : iteration 1970
the risk is 1923.43439948 : iteration 1971
the risk is 1923.40012972 : iteration 1972
the risk is 1923.36588459 : iteration 1973
the risk is 1923.33166404 : iteration 1974
the risk is 1923.29746805 : iteration 1975
the risk is 1923.2632966 : iteration 1976
the risk is 1923.22914966 : iteration 1977
the risk is 1923.1950272 : iteration 1978
the risk is 1923.16092919 : iteration 1979
the risk is 1923.1268556 : iteration 1980
the risk is 1923.09280641 : iteration 1981
the risk is 1923.05878159 : iteration 1982
the risk is 1923.02478111 : iteration 1983
the risk is 1922.99080495 : iteration 1984
the risk is 1922.95685308 : iteration 1985
the risk is 1922.92292547 : iteration 1986
the risk is 1922.88902209 : iteration 1987
the risk is 1922.85514292 : iteration 1988
the risk is 1922.82128793 : iteration 1989
the risk is 1922.78745709 : iteration 1990
the risk is 1922.75365038 : iteration 1991
the risk is 1922.71986777 : iteration 1992
the risk is 1922.68610923 : iteration 1993
the risk is 1922.65237474 : iteration 1994
the risk is 1922.61866426 : iteration 1995
the risk is 1922.58497778 : iteration 1996
the risk is 1922.55131526 : iteration 1997
the risk is 1922.51767668 : iteration 1998
the risk is 1922.48406202 : iteration 1999
the risk is 1922.45047124 : iteration 2000

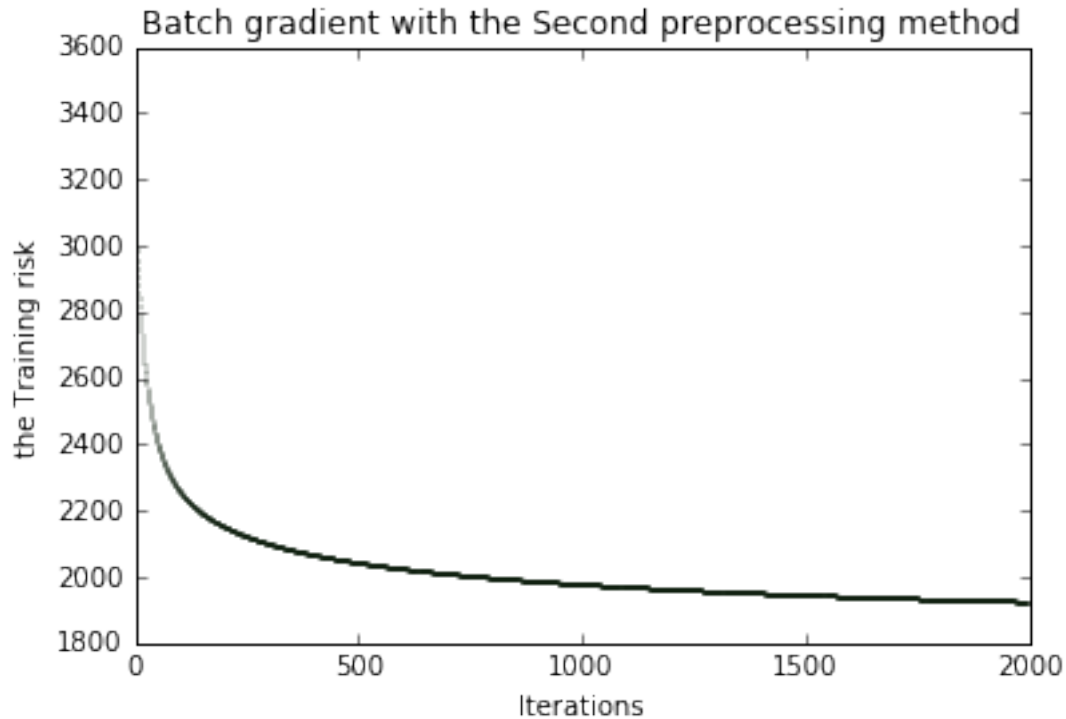
```

After 2000 iterations, the training risk decreases down to 1922.45047124, which is lower than what the first preprocessing method achieved. We plot the training risk over the number of iterations.

```

In [30]: xpos = np.arange(len(empirical_risk2[0]))+ 1
         plt.plot(xpos, empirical_risk2[0], 'go', ms = 0.4)
         plt.xlabel('Iterations')
         plt.ylabel('the Training risk')
         plt.title('Batch gradient with the Second preprocessing method')
         plt.show()

```



We run the batch gradient method with learning rate  $2 * 10^3$  for the third preprocessing data.

```
In [31]: empirical_risk3 = logistic_with_batch_gradient(w_0, sp_xtrain3, sp_ytrain, 2e-3, 2000)
```

```
the risk is 3584.95721786 : iteration 1
the risk is 4106.32058967 : iteration 2
the risk is 3197.9220317 : iteration 3
the risk is 3967.10472524 : iteration 4
the risk is 2682.77707914 : iteration 5
the risk is 2754.18057516 : iteration 6
the risk is 2463.95894291 : iteration 7
the risk is 2417.07294157 : iteration 8
the risk is 2271.68299745 : iteration 9
the risk is 2223.3775302 : iteration 10
the risk is 2155.00067837 : iteration 11
the risk is 2123.86911517 : iteration 12
the risk is 2089.30249473 : iteration 13
the risk is 2069.74877244 : iteration 14
the risk is 2050.14159893 : iteration 15
the risk is 2037.07824844 : iteration 16
the risk is 2024.64782528 : iteration 17
the risk is 2015.31136188 : iteration 18
the risk is 2006.66649052 : iteration 19
the risk is 1999.60404597 : iteration 20
the risk is 1993.12895398 : iteration 21
the risk is 1987.52734896 : iteration 22
the risk is 1982.38003115 : iteration 23
the risk is 1977.75244675 : iteration 24
```

```

the risk is 1820.99212377 : iteration 1969
the risk is 1820.99105226 : iteration 1970
the risk is 1820.98998186 : iteration 1971
the risk is 1820.98891256 : iteration 1972
the risk is 1820.98784437 : iteration 1973
the risk is 1820.98677729 : iteration 1974
the risk is 1820.98571113 : iteration 1975
the risk is 1820.98464642 : iteration 1976
the risk is 1820.98358263 : iteration 1977
the risk is 1820.98251994 : iteration 1978
the risk is 1820.98145834 : iteration 1979
the risk is 1820.98039784 : iteration 1980
the risk is 1820.97933843 : iteration 1981
the risk is 1820.97828011 : iteration 1982
the risk is 1820.97722287 : iteration 1983
the risk is 1820.97616673 : iteration 1984
the risk is 1820.97511166 : iteration 1985
the risk is 1820.97405768 : iteration 1986
the risk is 1820.97300478 : iteration 1987
the risk is 1820.97195296 : iteration 1988
the risk is 1820.97090222 : iteration 1989
the risk is 1820.96985255 : iteration 1990
the risk is 1820.96880396 : iteration 1991
the risk is 1820.96775644 : iteration 1992
the risk is 1820.96670999 : iteration 1993
the risk is 1820.96566461 : iteration 1994
the risk is 1820.9646203 : iteration 1995
the risk is 1820.96357705 : iteration 1996
the risk is 1820.96253487 : iteration 1997
the risk is 1820.96149375 : iteration 1998
the risk is 1820.96045369 : iteration 1999
the risk is 1820.95941468 : iteration 2000


```

After 2000 iterations, the training risk decreases down to 1820.95941468, which is the lowest among the three preprocessing method. We plot the training risk over the number of iterations.

```

In [32]: xpos = np.arange(len(empirical_risk3[0]))+ 1
         plt.plot(xpos, empirical_risk3[0], 'bo', ms =0.4)
         plt.xlabel('Iterations')
         plt.ylabel('the Training risk')
         plt.title('Batch gradient with the Third preprocessing method')
         plt.show()

```



### Stochastic Gradient Descent

In [3]: `import numpy.random as nr`

```
In [342]: def logistic_with_stochastic_gradient(w,X,y,e,n):
            w_temp = w
            risk = []
            nr.seed(0)
            indice = nr.choice(X.shape[0], n, replace =True)
            for i in np.arange(n):
                risk_temp = R(w_temp,X,y)
                j = indice[i]
                print ('the risk is ', risk_temp ,': iteration ', (i+1))
                w_temp = w_temp + e* (y[j]- s(np.dot(X[j,], w_temp)))* X[j,]
                risk = risk + [risk_temp]
            return risk
```

We run the stochastic gradient descent with learning rate  $10^{-2}$  on the first preprocessed data.

```
In [349]: training_risk1 = logistic_with_stochastic_gradient(w_0,sp_xtrain1,sp_ytrain,1e-2,3000)

the risk is  3584.95721786 : iteration  1
the risk is  3577.18569409 : iteration  2
the risk is  3569.54942435 : iteration  3
the risk is  3557.08974585 : iteration  4
the risk is  3549.43081846 : iteration  5
the risk is  3545.22638536 : iteration  6
the risk is  3544.27684629 : iteration  7
```

```

the risk is 2207.56654746 : iteration 2978
the risk is 2207.32699551 : iteration 2979
the risk is 2207.06153589 : iteration 2980
the risk is 2207.69725855 : iteration 2981
the risk is 2207.68775331 : iteration 2982
the risk is 2207.72057017 : iteration 2983
the risk is 2207.74726163 : iteration 2984
the risk is 2207.72799554 : iteration 2985
the risk is 2207.66483498 : iteration 2986
the risk is 2207.39742025 : iteration 2987
the risk is 2207.60431926 : iteration 2988
the risk is 2207.34799286 : iteration 2989
the risk is 2208.38634599 : iteration 2990
the risk is 2208.41171998 : iteration 2991
the risk is 2208.17888505 : iteration 2992
the risk is 2207.96174752 : iteration 2993
the risk is 2207.75378217 : iteration 2994
the risk is 2207.50733708 : iteration 2995
the risk is 2207.31335389 : iteration 2996
the risk is 2207.14605318 : iteration 2997
the risk is 2206.99784651 : iteration 2998
the risk is 2206.99700984 : iteration 2999
the risk is 2206.8646236 : iteration 3000

```

```
In [380]: np.amin(training_risk1)
```


```
Out[380]: 2201.0702252575302
```

The training risk decreases down to 2206.8646236 after 3000 iterations and the smallest risk we obtained is 2201.0702252575302. We plot the training risk over the number of iterations.

```

In [379]: xpos = np.arange(len(training_risk1))+ 1
          plt.plot(xpos, training_risk1, 'ro', ms =1)
          plt.xlabel('Iterations')
          plt.ylabel('the Training risk')
          plt.title('Stochastic gradient with the First preprocessing method')
          plt.show()

```



We run stochastic gradient descent on the second preprocessed data with learning rate  $6 * 10^{-3}$ .

```
In [374]: training_risk2 = logistic_with_stochastic_gradient(w_0,sp_xtrain2,sp_ytrain,6e-3,3000)
```

```
the risk is 3584.95721786 : iteration 1
the risk is 3212.15851585 : iteration 2
the risk is 3103.95591907 : iteration 3
the risk is 3100.14270579 : iteration 4
the risk is 3162.06581537 : iteration 5
the risk is 3254.15187477 : iteration 6
the risk is 3352.29449397 : iteration 7
the risk is 3062.88300343 : iteration 8
the risk is 3103.98918932 : iteration 9
the risk is 3166.53726792 : iteration 10
the risk is 3075.5836251 : iteration 11
the risk is 3040.81037176 : iteration 12
the risk is 3075.65130709 : iteration 13
the risk is 3152.06946332 : iteration 14
the risk is 3244.78102454 : iteration 15
the risk is 3335.65849333 : iteration 16
the risk is 3034.76868602 : iteration 17
the risk is 3063.65335317 : iteration 18
the risk is 3125.60909426 : iteration 19
the risk is 3203.31187791 : iteration 20
the risk is 3315.23172421 : iteration 21
the risk is 3420.23836054 : iteration 22
the risk is 3515.86146489 : iteration 23
the risk is 3607.81114933 : iteration 24
```

```

the risk is 2498.91498185 : iteration 2995
the risk is 2520.87270224 : iteration 2996
the risk is 2604.12963914 : iteration 2997
the risk is 2651.88789088 : iteration 2998
the risk is 2604.4552991 : iteration 2999
the risk is 2642.34550665 : iteration 3000

```

```
In [376]: np.amin(training_risk2)
```


```
Out[376]: 2080.0943692828496
```

After 3000 iteration, the minimum of the training risk we obtained is 2080.0943692828496. We plot the training risk over the number of iterations.

```

In [378]: xpos = np.arange(len(training_risk2))+ 1
plt.plot(xpos, training_risk2, 'go', ms =0.7)
plt.xlabel('Iterations')
plt.ylabel('the Training risk')
plt.title('Stochastic gradient with the Second preprocessing method')
plt.show()

```



We run the stochastic gradient descent with learning rate 0.25 on the third processed data.

```
In [390]: training_risk3 = logistic_with_stochastic_gradient(w_0,sp_xtrain3,sp_ytrain,0.25,3000)
```

```

the risk is 3584.95721786 : iteration 1
the risk is 3394.71136248 : iteration 2
the risk is 3263.56439499 : iteration 3

```



```

the risk is 2105.73797256 : iteration 2974
the risk is 2134.69381042 : iteration 2975
the risk is 2008.54715556 : iteration 2976
the risk is 1979.28252688 : iteration 2977
the risk is 2086.98264892 : iteration 2978
the risk is 2076.96616593 : iteration 2979
the risk is 2066.02879129 : iteration 2980
the risk is 2127.56851354 : iteration 2981
the risk is 2135.11036447 : iteration 2982
the risk is 2064.00367304 : iteration 2983
the risk is 2018.88532541 : iteration 2984
the risk is 2018.86484439 : iteration 2985
the risk is 2022.29227704 : iteration 2986
the risk is 1993.09866324 : iteration 2987
the risk is 1990.51389123 : iteration 2988
the risk is 1974.20998364 : iteration 2989
the risk is 1985.98964033 : iteration 2990
the risk is 1995.15919228 : iteration 2991
the risk is 2000.89313779 : iteration 2992
the risk is 2007.40622337 : iteration 2993
the risk is 2022.65532272 : iteration 2994
the risk is 2025.03936747 : iteration 2995
the risk is 2027.44919338 : iteration 2996
the risk is 2046.51990959 : iteration 2997
the risk is 2055.39792415 : iteration 2998
the risk is 2055.31713203 : iteration 2999
the risk is 2065.03700687 : iteration 3000

```

```
In [391]: np.amin(training_risk3)
```

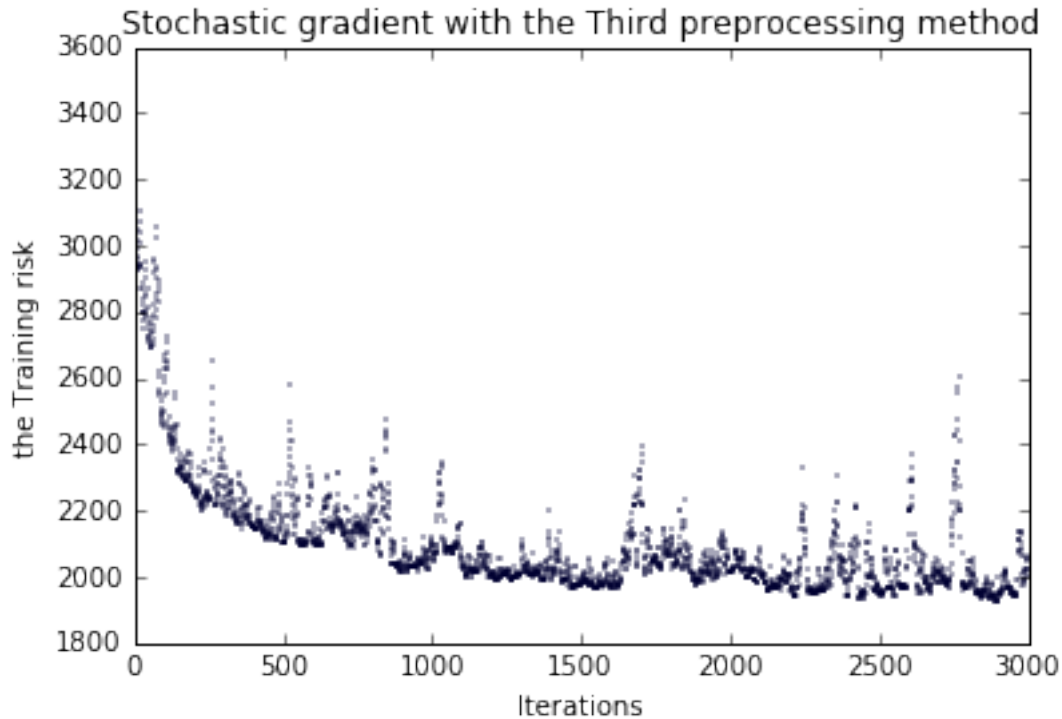
```
Out[391]: 1928.7904915537711
```

After 3000 iterations, the training risk decreases down to 2065.03700687 and the smallest risk we obtained is 1928.7904915537711. We plot the training risk over the number of iterations.

```

In [392]: xpos = np.arange(len(training_risk3))+ 1
          plt.plot(xpos, training_risk3, 'bo', ms =0.7)
          plt.xlabel('Iterations')
          plt.ylabel('the Training risk')
          plt.title('Stochastic gradient with the Third preprocessing method')
          plt.show()

```



**Stochastic gradient descent VS Batch gradient descent** It appears that the convergence rate, the speed in which a method converge to the minima, is higher in the batch gradient descent and the batch gradient method guarantees continuous decrease in the training risk. However, the stochastic gradient descent takes shorter time to compute for each step.

**Stochastic gradient descent with learning rate decreasing over iterations**

```
In [407]: def logistic_with_stochastic_gradient_decreasing_stepsize(w,X,y,e,n):
           w_temp = w
           risk = []
           nr.seed(0)
           indice = nr.choice(X.shape[0], n, replace =True)
           for i in np.arange(n):
               risk_temp = R(w_temp,X,y)
               j = indice[i]
               print ('the risk is ', risk_temp, ': iteration ', (i+1))
               w_temp = w_temp + (e/(i+1)) * (y[j]- s(np.dot(X[j,], w_temp)))* X[j,]
               risk = risk + [risk_temp]
           return risk
```

We run the stochastic gradient descent with learning rate  $\frac{1}{t}$  at step  $t$  on the first preprocessed data.

```
In [434]: training_risk_1 =logistic_with_stochastic_gradient_decreasing_stepsize(w_0,sp_xtrain1,sp_ytra

the risk is  3584.95721786 : iteration  1
the risk is  3252.40655651 : iteration  2
the risk is  3292.63860538 : iteration  3
the risk is  3085.89503634 : iteration  4
```

```

the risk is 2518.31359094 : iteration 2975
the risk is 2518.32229277 : iteration 2976
the risk is 2518.3229979 : iteration 2977
the risk is 2518.38436352 : iteration 2978
the risk is 2518.36244733 : iteration 2979
the risk is 2518.32153486 : iteration 2980
the risk is 2518.36607779 : iteration 2981
the risk is 2518.34297763 : iteration 2982
the risk is 2518.34759356 : iteration 2983
the risk is 2518.32727516 : iteration 2984
the risk is 2518.32661548 : iteration 2985
the risk is 2518.12750563 : iteration 2986
the risk is 2518.10939783 : iteration 2987
the risk is 2518.14073466 : iteration 2988
the risk is 2518.12265615 : iteration 2989
the risk is 2518.24819988 : iteration 2990
the risk is 2518.24551491 : iteration 2991
the risk is 2518.23155956 : iteration 2992
the risk is 2518.21762553 : iteration 2993
the risk is 2518.19963968 : iteration 2994
the risk is 2518.15902748 : iteration 2995
the risk is 2518.13735112 : iteration 2996
the risk is 2518.11944996 : iteration 2997
the risk is 2518.10149039 : iteration 2998
the risk is 2518.09542587 : iteration 2999
the risk is 2518.08163549 : iteration 3000


```

After 3000 iterations, the training risk decreases down to 2518.08163549. We plot the training risk over the number of iterations.

```

In [435]: xpos = np.arange(len(training_risk_1))+ 1
          plt.plot(xpos, training_risk_1, 'ro', ms =1)
          plt.xlabel('Iterations')
          plt.ylabel('the Training risk')
          plt.title('Stochastic gradient with decreasing rate on the First data')
          plt.show()

```



We run the stochastic gradient descent with the decreasing step size  $\frac{1.1}{t}$  at step  $t$ .

In [455]: `training_risk_2 =logistic_with_stochastic_gradient_decreasing_stepsize(w_0,sp_xtrain2,sp_ytra`

```

the risk is 3584.95721786 : iteration 1
the risk is 124707.212082 : iteration 2
the risk is 124707.212082 : iteration 3
the risk is 124707.212082 : iteration 4
the risk is 124707.212082 : iteration 5
the risk is 124707.212082 : iteration 6
the risk is 124707.212082 : iteration 7
the risk is 91134.397048 : iteration 8
the risk is 91134.397048 : iteration 9
the risk is 91134.397048 : iteration 10
the risk is 65899.9551046 : iteration 11
the risk is 65899.9551046 : iteration 12
the risk is 65899.9551046 : iteration 13
the risk is 65899.9551046 : iteration 14
the risk is 65899.9551046 : iteration 15
the risk is 65899.9551046 : iteration 16
the risk is 49957.5734774 : iteration 17
the risk is 49957.5734774 : iteration 18
the risk is 49957.5734774 : iteration 19
the risk is 49957.5734774 : iteration 20
the risk is 49957.5734774 : iteration 21
the risk is 49957.5734774 : iteration 22
the risk is 49957.5734774 : iteration 23
the risk is 49957.5734774 : iteration 24

```

```

the risk is 2421.23057353 : iteration 2995
the risk is 2418.74254936 : iteration 2996
the risk is 2415.20851453 : iteration 2997
the risk is 2412.41764359 : iteration 2998
the risk is 2412.59857607 : iteration 2999
the risk is 2411.97082038 : iteration 3000

```

```
In [457]: np.amin(training_risk_2)
```

```
Out[457]: 2404.8189851332536
```

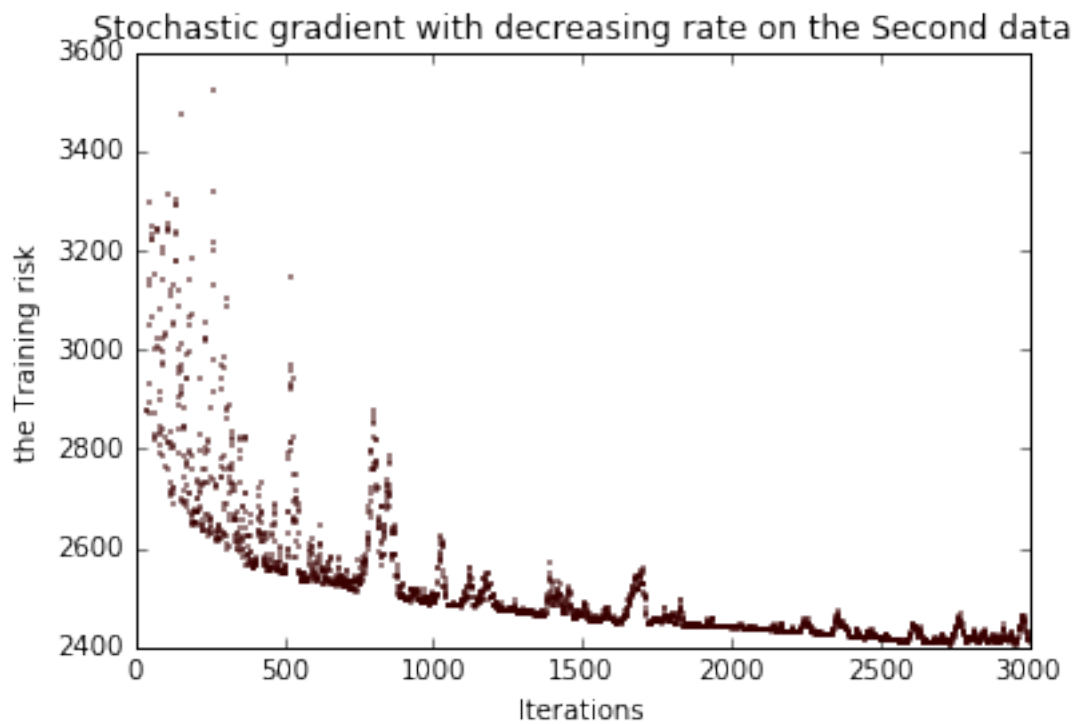
After 3000 iterations, the training risk decreases down to 2411.97082038 and the smallest risk we obtained is 2404.8189851332536. We plot the training risk over the number of iterations.

```

In [460]: xpos = np.arange(len(training_risk_2))+ 1

plt.plot(xpos, training_risk_2, 'ro', ms =1)
plt.xlabel('Iterations')
plt.ylabel('the Training risk')
plt.title('Stochastic gradient with decreasing rate on the Second data')
x1,x2,y1,y2 = plt.axis()
plt.axis((x1,x2,2400,3600))
plt.show()

```



We run the stochastic gradient descent with decreasing learning rate  $\frac{8}{t}$  at step  $t$  on the third preprocessed data.

```
In [473]: training_risk_3 =logistic_with_stochastic_gradient_decreasing_stepsize(w_0,sp_xtrain3,sp_ytra
```

the risk is 3584.95721786 : iteration 1  
the risk is 7630.11768245 : iteration 2  
the risk is 7632.57883964 : iteration 3  
the risk is 7634.56358585 : iteration 4  
the risk is 7686.37778253 : iteration 5  
the risk is 7754.55048236 : iteration 6  
the risk is 7755.60365454 : iteration 7  
the risk is 4867.83373023 : iteration 8  
the risk is 5108.06361153 : iteration 9  
the risk is 5109.29759209 : iteration 10  
the risk is 3928.76723755 : iteration 11  
the risk is 3993.73069164 : iteration 12  
the risk is 3998.25388454 : iteration 13  
the risk is 4046.87853921 : iteration 14  
the risk is 4056.59256502 : iteration 15  
the risk is 4057.69071289 : iteration 16  
the risk is 3638.55055255 : iteration 17  
the risk is 3639.86729894 : iteration 18  
the risk is 3675.36510593 : iteration 19  
the risk is 3676.41449269 : iteration 20  
the risk is 4122.11430793 : iteration 21  
the risk is 4200.39539916 : iteration 22  
the risk is 4232.40003714 : iteration 23  
the risk is 4262.45159182 : iteration 24  
the risk is 3597.43288372 : iteration 25  
the risk is 3630.37275309 : iteration 26  
the risk is 3416.90555764 : iteration 27  
the risk is 3231.67752292 : iteration 28  
the risk is 3067.06135222 : iteration 29  
the risk is 3068.01751595 : iteration 30  
the risk is 2837.21316484 : iteration 31  
the risk is 2861.54175094 : iteration 32  
the risk is 2882.57628875 : iteration 33  
the risk is 2884.26828219 : iteration 34  
the risk is 2828.53404188 : iteration 35  
the risk is 2796.81024109 : iteration 36  
the risk is 2797.53900777 : iteration 37  
the risk is 2732.19345743 : iteration 38  
the risk is 2732.61579409 : iteration 39  
the risk is 2731.71791174 : iteration 40  
the risk is 2731.69969785 : iteration 41  
the risk is 2697.96490653 : iteration 42  
the risk is 2701.75150852 : iteration 43  
the risk is 2691.7827796 : iteration 44  
the risk is 2685.52503714 : iteration 45  
the risk is 2682.24063911 : iteration 46  
the risk is 2682.9200891 : iteration 47  
the risk is 2677.89536187 : iteration 48  
the risk is 2682.55850291 : iteration 49  
the risk is 2679.22931728 : iteration 50  
the risk is 2670.30610646 : iteration 51  
the risk is 2671.0322023 : iteration 52  
the risk is 2671.6853807 : iteration 53  
the risk is 2671.95900081 : iteration 54

```

the risk is 2302.62544852 : iteration 2971
the risk is 2302.58945269 : iteration 2972
the risk is 2302.58648492 : iteration 2973
the risk is 2302.57391052 : iteration 2974
the risk is 2302.25005454 : iteration 2975
the risk is 2302.31946634 : iteration 2976
the risk is 2302.32903285 : iteration 2977
the risk is 2302.35492851 : iteration 2978
the risk is 2302.31174504 : iteration 2979
the risk is 2302.27500694 : iteration 2980
the risk is 2302.31435114 : iteration 2981
the risk is 2302.19095506 : iteration 2982
the risk is 2302.26699555 : iteration 2983
the risk is 2302.15816727 : iteration 2984
the risk is 2302.15528497 : iteration 2985
the risk is 2301.85693878 : iteration 2986
the risk is 2301.84030531 : iteration 2987
the risk is 2301.88186747 : iteration 2988
the risk is 2301.86694427 : iteration 2989
the risk is 2301.93227755 : iteration 2990
the risk is 2301.91759511 : iteration 2991
the risk is 2301.93009355 : iteration 2992
the risk is 2301.94260635 : iteration 2993
the risk is 2301.93036143 : iteration 2994
the risk is 2301.896843 : iteration 2995
the risk is 2301.85756132 : iteration 2996
the risk is 2301.84642162 : iteration 2997
the risk is 2301.82495288 : iteration 2998
the risk is 2301.82008546 : iteration 2999
the risk is 2301.83293867 : iteration 3000

```


After 3000 iterations, the training risk decreases down to 2301. We plot the training risk over the number of iterations.

```

In [474]: xpos = np.arange(len(training_risk_3))+ 1

plt.plot(xpos, training_risk_3, 'bo', ms =1)
plt.xlabel('Iterations')
plt.ylabel('the Training risk')
plt.title('Stochastic gradient with decreasing rate on the Third data')
x1,x2,y1,y2 = plt.axis()
plt.axis((x1,x2,2300,3600))
plt.show()

```



Is this strategy better than having a constant  $\epsilon$ ? With the constant step size (learning rate), the stochastic gradient produces lower training risk after the same number of iterations. This could be predictable in that the step size gets too small too fast when we use decreasing learning rates and therefore could not approach to the minimum as much as the stochastic gradient descent with constant step size does. It is worth to mention that the graph of training risk is much smoother curve with decreasing learning rate than with constant learning rate.

kernel logistic regression with a polynomial kernel of degree 2

```
In [19]: nr.seed(0)
         valid = nr.choice(sp_xtrain3.shape[0], sp_xtrain3.shape[0]//3, replace=False)
         sp_xvalid = sp_xtrain3[valid, ]
         sp_xtrain0 = sp_xtrain3[np.array([i not in valid for i in np.arange(sp_xtrain3.shape[0])]) ,:]
         sp_yvalid = sp_ytrain[valid]
         sp_ytrain0 = sp_ytrain[np.array([i not in valid for i in np.arange(sp_xtrain3.shape[0])])]
```

```
In [45]: sp_xtrain0.shape
         sp_xvalid.shape
```

```
Out[45]: (1724, 33)
```

```
In [12]: def quadratic_K(rho, X_train, X_test):
         return np.square(np.dot(X_train, np.transpose(X_test))) + rho)
```

```
In [13]: kf_ = KFold(3448, n_folds=10)
```

```
In [10]: # Risk function
         def R_(a, y_test, K):
```



```

temp = np.dot(K,a)
one = np.linspace(1,1, len(y_test))
return np.dot(np.transpose(1- y_test),temp) + np.dot(one, np.log(1+ np.exp(-temp)))

def stochastic_gradient_kernel(a,X,y, rho,e, lamda,n, K_):
    K__ = K_(rho,X, X)
    a_temp = a
    risk = []
    nr.seed(0)
    indice = nr.choice(X.shape[0], n, replace = True)
    for i in np.arange(n):
        # risk_temp = R_(a_temp, X, y, K)
        j = indice[i]
        # print ('the risk is ', risk_temp, ': iteration ', (i+1))
        b = lamda * a_temp
        b[j] = b[j] - y[j] + s(np.dot(K__[j,], a_temp))
        a_temp = a_temp - e * b
        # risk = risk + [risk_temp]
    return a_temp

def tenfold_cv_(rho, e, lamda, n, K_):
    R = 0
    for train_index, test_index in kf_:
        X_train, X_test = sp_xtrain0[train_index, :], sp_xtrain0[test_index, :]
        y_train, y_test = sp_ytrain0[train_index], sp_ytrain0[test_index]
        a0 = np.zeros(X_train.shape[0])
        a = stochastic_gradient_kernel(a0,X_train,y_train, rho,e, lamda,n, K_)
        temp = np.dot(a, np.square( np.dot(X_train, np.transpose(X_test))+ rho))
        one = np.linspace(1,1, len(y_test))
        R = R + np.dot(np.transpose(1- y_test),temp) + np.dot(one, np.log(1+ np.exp(-temp)))
    return R/10

```

In [94]: tenfold\_cv\_(0, 1e-5, 1e-3, 500, quadratic\_K)

Out[94]: 238.81979676288202

In [95]: tenfold\_cv\_(1e-3, 1e-5, 1e-3, 500, quadratic\_K)

Out[95]: 238.81964581654387

In [96]: tenfold\_cv\_(1e-1, 1e-5, 1e-3, 500, quadratic\_K)

Out[96]: 238.80433972893556

In [97]: tenfold\_cv\_(1e-2, 1e-5, 1e-3, 500, quadratic\_K)

Out[97]: 238.81828400128205

In [98]: tenfold\_cv\_(1, 1e-5, 1e-3, 500, quadratic\_K)

Out[98]: 238.6327606396122

In [99]: tenfold\_cv\_(10, 1e-5, 1e-3, 500, quadratic\_K)

Out[99]: 234.76740849818816

In [100]: tenfold\_cv\_(20, 1e-5, 1e-3, 500, quadratic\_K)

```
Out[100]: 232.1178878996912
```

```
In [101]: tenfold_cv_(30, 1e-5, 1e-3, 500, quadratic_K)
```

```
Out[101]: 235.93093307801436
```

```
In [102]: tenfold_cv_(18, 1e-5, 1e-3, 500, quadratic_K)
```

```
Out[102]: 232.13978447133462
```

```
In [103]: tenfold_cv_(19, 1e-5, 1e-3, 500, quadratic_K)
```

```
Out[103]: 232.09213791761448
```

10-fold cross validation shows that average training risk is the lowest when  $\rho = 19$ .

```
In [11]: def stochastic_gradient_kernel_print(X_train, X_test, y_train, y_test, rho, e, lamda, n, K_):
    K1 = K_(rho, X_train, X_train)
    K2 = K_(rho, X_train, X_test)
    risk = []
    nr.seed(0)
    indice = nr.choice(X_train.shape[0], n, replace = True)
    a0 = np.zeros(X_train.shape[0])
    temp = np.dot(a0, K2)
    one = np.linspace(1, 1, len(y_test))
    a_temp = a0
    for i in np.arange(n):
        risk_temp = np.dot(np.transpose(1- y_test), temp) + np.dot(one, np.log(1+ np.exp(-temp)))
        j = indice[i]
        print ('the risk is ', risk_temp, ': iteration ', (i+1))
        b = lamda * a_temp
        diff = y_train[j] - s(np.dot(a_temp, K1[j, ]))
        b[j] = b[j] - diff
        a_temp = a_temp - e * b
        temp = (1-e * lamda) * temp + e * diff * K2[j, ]
        risk = risk + [risk_temp]
    return [risk, a_temp]
```

```
In [87]: training_risk_kernel = stochastic_gradient_kernel_print(sp_xtrain0, sp_xtrain0,
                                                                sp_ytrain0, sp_ytrain0, 19, 1e-5, 1e-3,
```

```
the risk is 2389.97147857 : iteration 1
the risk is 2388.50723823 : iteration 2
the risk is 2386.91925433 : iteration 3
the risk is 2385.37438057 : iteration 4
the risk is 2383.98065696 : iteration 5
the risk is 2385.3855362 : iteration 6
the risk is 2386.7790168 : iteration 7
the risk is 2385.25283763 : iteration 8
the risk is 2383.78251136 : iteration 9
the risk is 2382.38667382 : iteration 10
the risk is 2383.79471226 : iteration 11
the risk is 2382.27707668 : iteration 12
the risk is 2380.88604362 : iteration 13
the risk is 2379.4186275 : iteration 14
the risk is 2378.00834721 : iteration 15
```

```

the risk is 1733.23091165 : iteration 29986
the risk is 1733.34676539 : iteration 29987
the risk is 1733.25907708 : iteration 29988
the risk is 1733.17172024 : iteration 29989
the risk is 1733.08556903 : iteration 29990
the risk is 1733.27284225 : iteration 29991
the risk is 1733.19904327 : iteration 29992
the risk is 1733.130197 : iteration 29993
the risk is 1733.08688898 : iteration 29994
the risk is 1733.31562529 : iteration 29995
the risk is 1733.26332413 : iteration 29996
the risk is 1733.35728836 : iteration 29997
the risk is 1733.57522346 : iteration 29998
the risk is 1733.48289971 : iteration 29999
the risk is 1733.59507488 : iteration 30000

```

After the 30000 iteration, we achieve the lowest training risk 1733.59507488. We plot the training risk over the number of iterations.

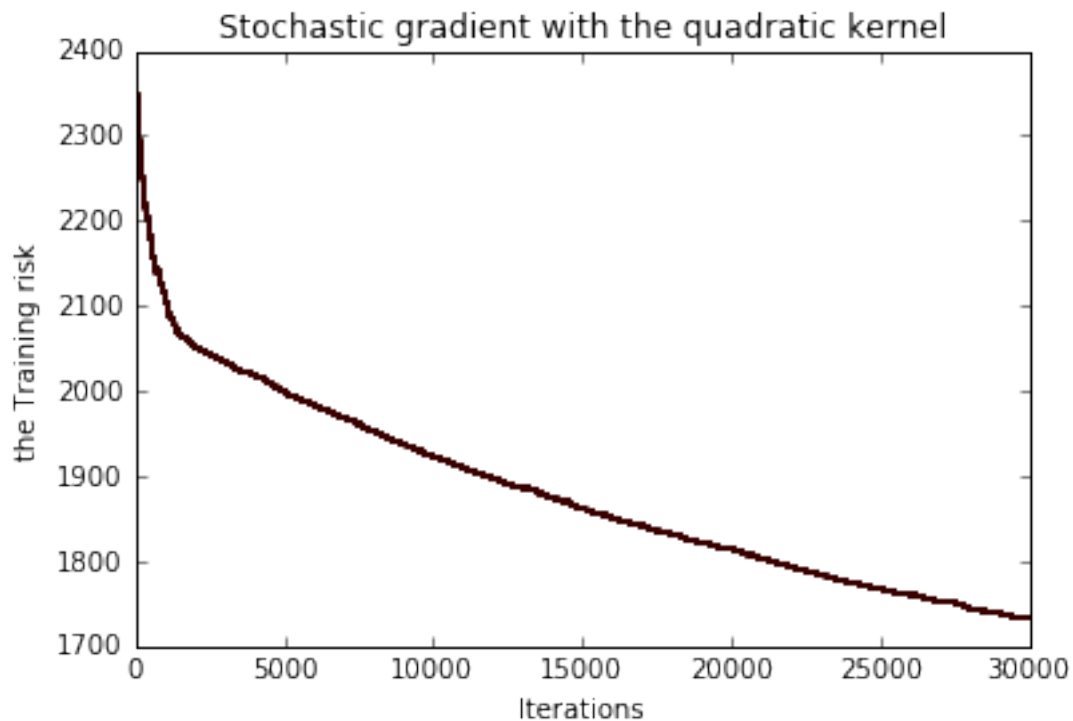
```

In [88]: xpos = np.arange(len(training_risk_kernel[0]))+ 1

plt.plot(xpos, training_risk_kernel[0], 'ro', ms =1)
plt.xlabel('Iterations')
plt.ylabel('the Training risk')
plt.title('Stochastic gradient with the quadratic kernel')

plt.show()

```



```
In [90]: validation_risk_kernel = stochastic_gradient_kernel_print(sp_xtrain0, sp_xvalid, sp_ytrain0,
                                                                    sp_yvalid, 19, 1e-5, 1e-3, 30000, quad)
```

```
the risk is 1194.98573929 : iteration 1
the risk is 1194.20773155 : iteration 2
the risk is 1193.36570734 : iteration 3
the risk is 1192.55076004 : iteration 4
the risk is 1191.80927024 : iteration 5
the risk is 1192.55777194 : iteration 6
the risk is 1193.30053708 : iteration 7
the risk is 1192.48966295 : iteration 8
the risk is 1191.70459066 : iteration 9
the risk is 1190.96184941 : iteration 10
the risk is 1191.7099116 : iteration 11
the risk is 1190.90356364 : iteration 12
the risk is 1190.16327242 : iteration 13
the risk is 1189.3846922 : iteration 14
the risk is 1188.63467339 : iteration 15
the risk is 1187.88011407 : iteration 16
the risk is 1187.10411683 : iteration 17
the risk is 1187.89314283 : iteration 18
the risk is 1187.09505982 : iteration 19
the risk is 1186.31956597 : iteration 20
the risk is 1185.5250791 : iteration 21
the risk is 1184.80206725 : iteration 22
the risk is 1184.07914126 : iteration 23
the risk is 1184.81928945 : iteration 24
the risk is 1184.05880372 : iteration 25
the risk is 1183.34059455 : iteration 26
the risk is 1182.58525699 : iteration 27
the risk is 1181.87188959 : iteration 28
the risk is 1181.06952257 : iteration 29
the risk is 1181.80351573 : iteration 30
the risk is 1182.53824473 : iteration 31
the risk is 1181.82484125 : iteration 32
the risk is 1181.11381816 : iteration 33
the risk is 1180.37291976 : iteration 34
the risk is 1179.66487554 : iteration 35
the risk is 1180.42071671 : iteration 36
the risk is 1179.70344952 : iteration 37
the risk is 1178.91792199 : iteration 38
the risk is 1179.64837315 : iteration 39
the risk is 1180.45956979 : iteration 40
the risk is 1181.19454139 : iteration 41
the risk is 1180.4698393 : iteration 42
the risk is 1179.67675291 : iteration 43
the risk is 1180.40076674 : iteration 44
the risk is 1179.62847651 : iteration 45
the risk is 1178.9219093 : iteration 46
the risk is 1178.09916193 : iteration 47
the risk is 1178.85748216 : iteration 48
the risk is 1179.61947512 : iteration 49
the risk is 1178.90099568 : iteration 50
the risk is 1179.69532893 : iteration 51
```

```

the risk is 849.327567256 : iteration 29968
the risk is 849.287251117 : iteration 29969
the risk is 849.243460775 : iteration 29970
the risk is 849.180511968 : iteration 29971
the risk is 849.101660624 : iteration 29972
the risk is 849.029430679 : iteration 29973
the risk is 848.937254226 : iteration 29974
the risk is 848.878913036 : iteration 29975
the risk is 848.820981821 : iteration 29976
the risk is 848.753475218 : iteration 29977
the risk is 848.702492902 : iteration 29978
the risk is 848.792638151 : iteration 29979
the risk is 848.708065881 : iteration 29980
the risk is 848.641003732 : iteration 29981
the risk is 848.608099105 : iteration 29982
the risk is 848.729436766 : iteration 29983
the risk is 848.64519449 : iteration 29984
the risk is 848.599187664 : iteration 29985
the risk is 848.540981381 : iteration 29986
the risk is 848.636837815 : iteration 29987
the risk is 848.570570446 : iteration 29988
the risk is 848.504476826 : iteration 29989
the risk is 848.439001153 : iteration 29990
the risk is 848.580956226 : iteration 29991
the risk is 848.528980923 : iteration 29992
the risk is 848.479532176 : iteration 29993
the risk is 848.450313802 : iteration 29994
the risk is 848.618014203 : iteration 29995
the risk is 848.58231981 : iteration 29996
the risk is 848.662712355 : iteration 29997
the risk is 848.825113273 : iteration 29998
the risk is 848.756474267 : iteration 29999
the risk is 848.848799934 : iteration 30000

```

The validation risk, which initially was 1194.9857, decreases down to 848.848799934 after 30000 iterations.


```

In [91]: xpos = np.arange(len(validation_risk_kernel[0]))+ 1

plt.plot(xpos, validation_risk_kernel[0], 'go', ms =1)
plt.xlabel('Iterations')
plt.ylabel('the Validation risk')
plt.title('Stochastic gradient with the quadratic kernel')

plt.show()

```



Repeat the same experiment with the linear kernel

In [92]: *# define the linear kernel function*

```
def linear_K(rho, X_train, X_test):
    return np.dot(X_train, np.transpose(X_test)) + rho
```

10-fold Cross Validation

In [104]: `tenfold_cv_(0, 1e-5, 1e-3, 500, linear_K)`

Out[104]: 238.81973628796601

In [105]: `tenfold_cv_(1, 1e-5, 1e-3, 500, linear_K)`

Out[105]: 238.63221702134652

In [106]: `tenfold_cv_(20, 1e-5, 1e-3, 500, linear_K)`

Out[106]: 232.59098370465694

In [107]: `tenfold_cv_(10, 1e-5, 1e-3, 500, linear_K)`

Out[107]: 234.53067277148665

In [108]: `tenfold_cv_(30, 1e-5, 1e-3, 500, linear_K)`

Out[108]: 252.04620848751625

In [109]: `tenfold_cv_(19, 1e-5, 1e-3, 500, linear_K)`

```
Out[109]: 232.18632045192498
```

```
In [110]: tenfold_cv_(18, 1e-5, 1e-3, 500, linear_K)
```

```
Out[110]: 231.97186548732543
```

```
In [111]: tenfold_cv_(15, 1e-5, 1e-3, 500, linear_K)
```

```
Out[111]: 232.2471268876871
```

```
In [112]: tenfold_cv_(16, 1e-5, 1e-3, 500, linear_K)
```

```
Out[112]: 232.02402140148493
```

```
In [113]: tenfold_cv_(17, 1e-5, 1e-3, 500, linear_K)
```

```
Out[113]: 231.92511545177331
```

It is when  $\rho = 17$  that the 10-fold cross validation records the lowest average risk.

```
In [116]: training_risk_kernel_linear = stochastic_gradient_kernel_print(sp_xtrain0, sp_xtrain0,  
                                                                           sp_ytrain0, sp_ytrain0, 17, 1e-5, 1e-3
```

```
the risk is 2389.97147857 : iteration 1  
the risk is 2389.90650462 : iteration 2  
the risk is 2389.8385066 : iteration 3  
the risk is 2389.77134939 : iteration 4  
the risk is 2389.70773888 : iteration 5  
the risk is 2389.77139251 : iteration 6  
the risk is 2389.83472035 : iteration 7  
the risk is 2389.76799327 : iteration 8  
the risk is 2389.70251425 : iteration 9  
the risk is 2389.63874838 : iteration 10  
the risk is 2389.70253505 : iteration 11  
the risk is 2389.6357788 : iteration 12  
the risk is 2389.57202282 : iteration 13  
the risk is 2389.50624651 : iteration 14  
the risk is 2389.4417862 : iteration 15  
the risk is 2389.37710109 : iteration 16  
the risk is 2389.31116399 : iteration 17  
the risk is 2389.37715804 : iteration 18  
the risk is 2389.3100696 : iteration 19  
the risk is 2389.24373855 : iteration 20  
the risk is 2389.17696497 : iteration 21  
the risk is 2389.11327866 : iteration 22  
the risk is 2389.04945504 : iteration 23  
the risk is 2389.1132089 : iteration 24  
the risk is 2389.04750431 : iteration 25  
the risk is 2388.98383783 : iteration 26  
the risk is 2388.91815372 : iteration 27  
the risk is 2388.85450706 : iteration 28  
the risk is 2388.78653727 : iteration 29  
the risk is 2388.85018377 : iteration 30  
the risk is 2388.91378182 : iteration 31  
the risk is 2388.85013542 : iteration 32  
the risk is 2388.78649886 : iteration 33
```

the risk is 2099.46398546 : iteration 29950  
the risk is 2099.45182715 : iteration 29951  
the risk is 2099.43939032 : iteration 29952  
the risk is 2099.46167567 : iteration 29953  
the risk is 2099.48244462 : iteration 29954  
the risk is 2099.469774 : iteration 29955  
the risk is 2099.45665782 : iteration 29956  
the risk is 2099.44473125 : iteration 29957  
the risk is 2099.43191747 : iteration 29958  
the risk is 2099.45189557 : iteration 29959  
the risk is 2099.47233093 : iteration 29960  
the risk is 2099.46016775 : iteration 29961  
the risk is 2099.4475425 : iteration 29962  
the risk is 2099.43506049 : iteration 29963  
the risk is 2099.45804901 : iteration 29964  
the risk is 2099.48040111 : iteration 29965  
the risk is 2099.46771892 : iteration 29966  
the risk is 2099.48816267 : iteration 29967  
the risk is 2099.47569469 : iteration 29968  
the risk is 2099.46231486 : iteration 29969  
the risk is 2099.44914244 : iteration 29970  
the risk is 2099.4364409 : iteration 29971  
the risk is 2099.42442745 : iteration 29972  
the risk is 2099.41223704 : iteration 29973  
the risk is 2099.40031315 : iteration 29974  
the risk is 2099.38766452 : iteration 29975  
the risk is 2099.37501964 : iteration 29976  
the risk is 2099.36285995 : iteration 29977  
the risk is 2099.34991765 : iteration 29978  
the risk is 2099.37047535 : iteration 29979  
the risk is 2099.35866724 : iteration 29980  
the risk is 2099.34653594 : iteration 29981  
the risk is 2099.33298303 : iteration 29982  
the risk is 2099.35505313 : iteration 29983  
the risk is 2099.34294781 : iteration 29984  
the risk is 2099.32999438 : iteration 29985  
the risk is 2099.31740012 : iteration 29986  
the risk is 2099.33811277 : iteration 29987  
the risk is 2099.32590296 : iteration 29988  
the risk is 2099.31378006 : iteration 29989  
the risk is 2099.30166082 : iteration 29990  
the risk is 2099.32389443 : iteration 29991  
the risk is 2099.31120034 : iteration 29992  
the risk is 2099.29829192 : iteration 29993  
the risk is 2099.28467202 : iteration 29994  
the risk is 2099.30780137 : iteration 29995  
the risk is 2099.29442467 : iteration 29996  
the risk is 2099.31464708 : iteration 29997  
the risk is 2099.33758979 : iteration 29998  
the risk is 2099.32546177 : iteration 29999  
the risk is 2099.345999 : iteration 30000

After 30000 iterations, the training risk decreases down to 2099.345999 in logistic ridge regression with




the linear kernel. We plot the training risk over the number of iterations.

```
In [119]: xpos = np.arange(len(training_risk_kernel_linear[0]))+ 1

plt.plot(xpos, training_risk_kernel_linear[0], 'ro', ms =1)
plt.xlabel('Iterations')
plt.ylabel('the Training risk')
plt.title('Stochastic gradient with the linear kernel')

plt.show()
```



### Validation risk

```
In [118]: validation_risk_kernel_linear = stochastic_gradient_kernel_print(sp_xtrain0, sp_xvalid,
                                                                           sp_ytrain0, sp_yvalid, 19, 1e-5, 1e-3,
```

```
the risk is 1194.98573929 : iteration 1
the risk is 1194.94743648 : iteration 2
the risk is 1194.90757046 : iteration 3
the risk is 1194.86825616 : iteration 4
the risk is 1194.83066519 : iteration 5
the risk is 1194.86830714 : iteration 6
the risk is 1194.90578659 : iteration 7
the risk is 1194.86656239 : iteration 8
the risk is 1194.82790979 : iteration 9
the risk is 1194.7902374 : iteration 10
the risk is 1194.82789426 : iteration 11
```

```

the risk is 1032.3747848 : iteration 29982
the risk is 1032.38861913 : iteration 29983
the risk is 1032.38133567 : iteration 29984
the risk is 1032.37357353 : iteration 29985
the risk is 1032.36602824 : iteration 29986
the risk is 1032.37914496 : iteration 29987
the risk is 1032.37173782 : iteration 29988
the risk is 1032.36439965 : iteration 29989
the risk is 1032.35706368 : iteration 29990
the risk is 1032.37094066 : iteration 29991
the risk is 1032.36338388 : iteration 29992
the risk is 1032.35567309 : iteration 29993
the risk is 1032.34763211 : iteration 29994
the risk is 1032.36201619 : iteration 29995
the risk is 1032.35408391 : iteration 29996
the risk is 1032.36694375 : iteration 29997
the risk is 1032.38125155 : iteration 29998
the risk is 1032.3739098 : iteration 29999
the risk is 1032.38697496 : iteration 30000

```

The validation risk, which initially was 1194.9857, decreases down to 1032.38697. We plot the validation risk over the number of iterations.

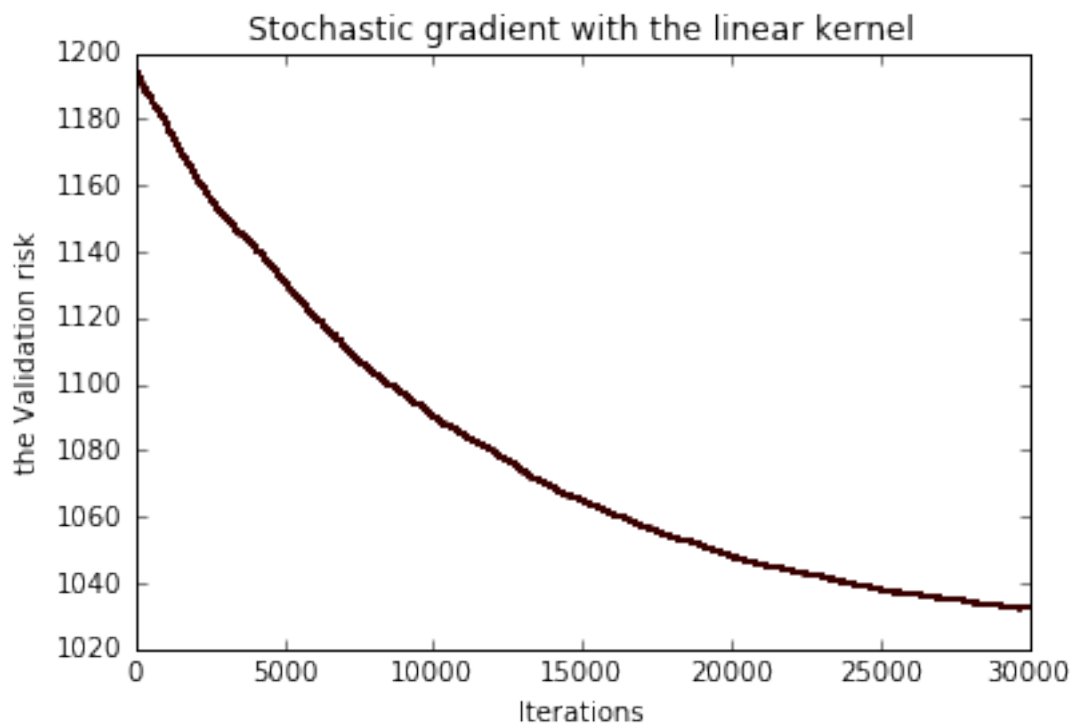
```

In [121]: xpos = np.arange(len(validation_risk_kernel_linear[0]))+ 1

plt.plot(xpos, validation_risk_kernel_linear[0], 'ro', ms =1)
plt.xlabel('Iterations')
plt.ylabel('the Validation risk')
plt.title('Stochastic gradient with the linear kernel')

plt.show()

```



**Does the quadratic kernel overfit the data?** According to the validation risk, the quadratic kernel does not appear to overfit the data. It gives lower validation risk than the linear kernel and the proportions of numbers of data and the total risk are almost the same in the validation risk and the training risk.

**For each kernel, should you decrease or increase  $\lambda$  to try to improve performance?**

```
In [123]: tenfold_cv_(17, 1e-5, 1e-3, 500, linear_K)
```

```
Out[123]: 231.92511545177331
```

```
In [126]: tenfold_cv_(17, 1e-5, 1e-2, 500, linear_K)
```

```
Out[126]: 231.92515041100859
```

```
In [128]: tenfold_cv_(17, 1e-5, 1, 500, linear_K)
```

```
Out[128]: 231.92902977056451
```

```
In [130]: tenfold_cv_(17, 1e-5, 10, 500, linear_K)
```

```
Out[130]: 231.96726038354353
```

```
In [132]: tenfold_cv_(17, 1e-5, 1e-5, 500, linear_K)
```

```
Out[132]: 231.92511160659714
```

```
In [133]: tenfold_cv_(19, 1e-5, 1e-5, 500, quadratic_K)
```

```
Out[133]: 232.09213334165784
```

```
In [134]: tenfold_cv_(19, 1e-5, 1e-3, 500, quadratic_K)
```

```
Out[134]: 232.09213791761448
```

```
In [135]: tenfold_cv_(19, 1e-5, 1e-1, 500, quadratic_K)
```

```
Out[135]: 232.09259575993278
```

```
In [137]: tenfold_cv_(19, 1e-5, 10, 500, quadratic_K)
```

```
Out[137]: 232.14075190405794
```

```
In [138]: tenfold_cv_(19, 1e-5, 100, 500, quadratic_K)
```

```
Out[138]: 232.72369887001486
```

We tried several different  $\lambda$  values in cross validation to see if they can improve performance. There was no noticeable decrease in the training risk compared to the training risk when  $\lambda$  is  $10^{-3}$ . This implies that there is no need to decrease or increase  $\lambda$  to improve performance.

**Kaggle** Based on training risks, or empirical risks, of each classifier, the best classifier is the logistic regression with batch gradient descent along with the third preprocessed data. We submitted the predicted labels to Kaggle that is derived by the best classifier, and the kaggle score was 0.78231.

```
In [35]: empirical_risk = logistic_with_batch_gradient(w_0, sp_xtrain3, sp_ytrain, 2e-3, 20000)
```

```
the risk is 3584.95721786 : iteration 1
the risk is 4106.32058967 : iteration 2
the risk is 3197.9220317 : iteration 3
the risk is 3967.10472524 : iteration 4
the risk is 2682.77707914 : iteration 5
the risk is 2754.18057516 : iteration 6
the risk is 2463.95894291 : iteration 7
the risk is 2417.07294157 : iteration 8
the risk is 2271.68299745 : iteration 9
the risk is 2223.3775302 : iteration 10
the risk is 2155.00067837 : iteration 11
the risk is 2123.86911517 : iteration 12
the risk is 2089.30249473 : iteration 13
the risk is 2069.74877244 : iteration 14
the risk is 2050.14159893 : iteration 15
the risk is 2037.07824844 : iteration 16
the risk is 2024.64782528 : iteration 17
the risk is 2015.31136188 : iteration 18
the risk is 2006.66649052 : iteration 19
the risk is 1999.60404597 : iteration 20
the risk is 1993.12895398 : iteration 21
the risk is 1987.52734896 : iteration 22
the risk is 1982.38003115 : iteration 23
the risk is 1977.75244675 : iteration 24
the risk is 1973.46265985 : iteration 25
the risk is 1969.50619858 : iteration 26
the risk is 1965.79984485 : iteration 27
the risk is 1962.32388258 : iteration 28
the risk is 1959.03723767 : iteration 29
the risk is 1955.92161228 : iteration 30
```

```

the risk is 1819.05681309 : iteration 19957
the risk is 1819.05680154 : iteration 19958
the risk is 1819.05678999 : iteration 19959
the risk is 1819.05677845 : iteration 19960
the risk is 1819.0567669 : iteration 19961
the risk is 1819.05675536 : iteration 19962
the risk is 1819.05674381 : iteration 19963
the risk is 1819.05673227 : iteration 19964
the risk is 1819.05672073 : iteration 19965
the risk is 1819.05670919 : iteration 19966
the risk is 1819.05669765 : iteration 19967
the risk is 1819.05668611 : iteration 19968
the risk is 1819.05667458 : iteration 19969
the risk is 1819.05666304 : iteration 19970
the risk is 1819.05665151 : iteration 19971
the risk is 1819.05663997 : iteration 19972
the risk is 1819.05662844 : iteration 19973
the risk is 1819.05661691 : iteration 19974
the risk is 1819.05660538 : iteration 19975
the risk is 1819.05659385 : iteration 19976
the risk is 1819.05658233 : iteration 19977
the risk is 1819.0565708 : iteration 19978
the risk is 1819.05655927 : iteration 19979
the risk is 1819.05654775 : iteration 19980
the risk is 1819.05653623 : iteration 19981
the risk is 1819.0565247 : iteration 19982
the risk is 1819.05651318 : iteration 19983
the risk is 1819.05650166 : iteration 19984
the risk is 1819.05649015 : iteration 19985
the risk is 1819.05647863 : iteration 19986
the risk is 1819.05646711 : iteration 19987
the risk is 1819.0564556 : iteration 19988
the risk is 1819.05644408 : iteration 19989
the risk is 1819.05643257 : iteration 19990
the risk is 1819.05642106 : iteration 19991
the risk is 1819.05640955 : iteration 19992
the risk is 1819.05639804 : iteration 19993
the risk is 1819.05638653 : iteration 19994
the risk is 1819.05637502 : iteration 19995
the risk is 1819.05636351 : iteration 19996
the risk is 1819.05635201 : iteration 19997
the risk is 1819.05634051 : iteration 19998
the risk is 1819.056329 : iteration 19999
the risk is 1819.0563175 : iteration 20000

```

```

In [41]: temp = np.dot(sp_test, empirical_risk[1])
        prob = s(temp)
        sp_ytest = prob >= 1/2
        sp_ypred = np.asarray([[i+1, sp_ytest[i]] for i in np.arange(5857)])
        np.savetxt('sp_ytest2.csv', sp_ypred, fmt = '%1.1u' , delimiter = ',', header = 'Id,Category',c

```

### 1.3.1 Problem 4: Revisiting Logistic Regression

Recall that in logistic regression, the logistic function,  $s(z) = \frac{1}{1+e^{-z}}$ , is used to map output values between  $[0, 1]$ . Consider the function  $g(z) = \frac{\tanh(z)+1}{2}$ , where  $\tanh(z) = 2s(2z) - 1$ , that also maps outputs between  $[0, 1]$ . In this problem, we will explore using  $g(z)$  instead of the logistic function in logistic regression.

$$g(z) = \frac{\tanh(z) + 1}{2} = \frac{2s(2z)}{2} = \frac{2}{2 + 2e^{-2z}} = \frac{1}{2} + \frac{1 - e^{-2z}}{2 + 2e^{-2z}} = \frac{1}{2} + \frac{e^z - e^{-z}}{2(e^z + e^{-z})}.$$

and

$$g'(z) = \frac{dg(z)}{dz} = \frac{2(e^z + e^{-z})(e^z + e^{-z}) - 2(e^z - e^{-z})^2}{4(e^z + e^{-z})^2} = \frac{1}{2} - \frac{\tanh(z)^2}{2}$$

The appropriate batch gradient ascent update function is

$$\begin{aligned} w &\leftarrow w - \lambda \nabla J(w) \\ &= w - \lambda \sum_{i=1}^n \left( \frac{y_i g'(X_i \cdot w)}{g(X_i \cdot w)} - \frac{(1 - y_i) g'(X_i \cdot w)}{1 - g(X_i \cdot w)} \right) X_i \\ &= w - \lambda \sum_{i=1}^n \left( \frac{y_i \left( \frac{1}{2} - \frac{\tanh^2(X_i \cdot w)}{2} \right)}{\frac{\tanh(X_i \cdot w) + 1}{2}} - \frac{(1 - y_i) \left( \frac{1}{2} - \frac{\tanh^2(X_i \cdot w)}{2} \right)}{1 - \frac{\tanh(X_i \cdot w) + 1}{2}} \right) X_i \\ &= w - \lambda \sum_{i=1}^n (2y_i - 1 - \tanh(X_i \cdot w)) X_i \end{aligned}$$

### 1.3.2 Problem 5: Real World Spam Classification

Linear SVM can not utilize the new feature as he desired, because the spams usually are sent around midnight and number of milliseconds since the previous midnight are concentrated in either  $(0, t]$  or  $[86400000 - t, 86400000)$  for small  $t$  and linear SVM can not classify whether this new feature is either in one of the two intervals  $(0, t]$  and  $[86400000 - t, 86400000)$  or not. Therefore in order to classify the feature correctly with linear SVM, Daniel may want to try one of following methods: 1. introduce the new feature as the number of milliseconds since the previous noon, rather than since the previous midnight so that the time around midnight could have continuous possible values, 2. use quadratic kernel instead of linear kernel so that the linear SVM can classify whether the new feature is either in one of the two intervals  $(0, t]$  and  $[86400000 - t, 86400000)$  or not.